

FOMO to Full Proactive and Autonomous Brevio Map

Timeline and co-founder operating map

Purpose: This document keeps the full Brevio vision visible while the team moves fast through FOMO. It is a co-founder operating map for the transition from FOMO wedge to proactive, autonomous, memory-rich personal agent.

FOMO is the runway. Brevio is the full proactive and autonomous personal agent.

1. North Star

Brevio should become a human-feeling, iMessage-native personal AI assistant that remembers the user, learns preferences, understands goals, chooses tools, asks for missing access, uses tools safely, and eventually completes real-world tasks with approval.

Text this number. Brevio knows me. It helps me. It asks before doing anything risky.

Under the hood, Brevio must behave like a governed agent operating system: secure OAuth, permission gates, audit logs, state machines, smoke tests, feedback/memory loops, idempotency, observability, and production-grade reliability.

2. Honest Current State

Current Brevio/FOMO is not yet the full proactive autonomous agent. It is the first wedge and kernel slice.

- What exists: Gmail-first FOMO workflow, Slack review, SendBlue/iMessage rail, reply parser, STOP/START, memory/feedback foundation, HMR emerging, reliability hardening, private onboarding, per-user isolation.
- What is missing: broad proactive learning, full personal context engine, calendar awareness, coffee/stocks/vacation routines, approved actions, browser automation fallback, multimodal perception, delegated agents, production scale.

This is acceptable only if each FOMO phase builds or validates reusable Brevio agent layers.

3. Why We Do Not Jump Unsafely To Full Autonomy

Full proactive Brevio would observe private life context, form memories, propose help, and sometimes act through tools. Without the right layers, it can become wrong, creepy, or dangerous.

Brevio can learn patterns. Brevio can propose help. Brevio cannot silently act on high-risk patterns.

Scenario	Correct Brevio Behavior	Unsafe Behavior
Coffee	Notice routine, ask whether to remind or prepare an order for approval.	Silently purchase coffee.
Stocks	Learn watch interests and offer alerts with permission.	Trade silently or make reckless financial actions.
Vacation	Ask if the user wants trip organization help after signals appear.	Silently book flights/hotels or make purchases.
Calendar	Use read-only schedule context first; ask before writes.	Add/delete events without exact approval.

4. Core Brevio Agent Dimensions

Dimension	Long-term meaning	Current FOMO status	Proof of progress
Autonomy	Reason about needs, missing tools, and safe next steps.	Limited to FOMO ranking/replies/gates.	Typed plans, approvals, audit, tests.
Proactivity	Notice patterns/events before the user asks.	Gmail alerting only.	Context-triggered suggestions with consent.
Memory Architecture	Source-linked, per-user, correctable/deletable memories.	Basic signals/events.	Memory has source/confidence/recency /user_id.
Agent Core + Reasoning	Ranking, planning, uncertainty, tool choice.	Classifiers and reasons.	Evals, schema validation, smoke behavior.
Tool / Workflow Orchestration	Governed tools and multi-step workflows.	Gmail/Slack/SendBlue workflow.	Registry, states, retries, audit.
Security / Permission Gates	AI proposes; system gates; user approves.	Strong current layer.	Kill switches, egress, policy, approvals.
Multimodal + Perception	Future files/images/screens/docs perception.	Not active.	Provider-specific eval/smoke gates.
Feedback + Learn/Grow	Natural feedback improves future behavior.	Partial; needs substrate.	Feedback -> memory -> behavior change.
Human Message Renderer	Context becomes human help.	Email alert surface emerging.	Tone/length/leak tests.
Observability / Evals / Reliability	Debuggable, measurable, recoverable agent.	Strong for FOMO smoke path.	Evidence scripts, runbooks, regressions.
Production Scale	Hundreds/thousands safely.	Not active.	Queues, monitoring, support, compliance.
User Trust / Consent	User controls what Brevio reads/remembers/does.	Strong in friend beta.	Privacy copy, deletion, STOP, approvals.

5. FOMO As Runway

FOMO proves the loop: context -> reasoning -> approval -> action -> reply -> learning -> audit.

FOMO piece	Reusable Brevio layer
Gmail polling	Context provider pattern
Ranker	Agent reasoning / importance model
Slack approval	Human approval checkpoint

SendBlue/iMessage	Messaging/action channel
STOP/START	Consent/control semantics
Reply parser	Natural feedback input
Memory signals	Personal memory architecture
Human Message Renderer	Universal product voice layer
Smoke tests	Real-world provider proof
Audit logs	Agent observability

Rule: FOMO-only polish should be deferred unless it fixes a real user trust blocker or builds a reusable Brevio layer.

6. Timeline And Phase Plan

Estimates assume Claude-heavy build cycles, quick founder review, and provider availability. They are planning estimates, not promises.

Phase	Goal	Fast pace	Normal pace	Deliverable	Proof gate
A. Finish immediate HMR runway	Complete current HMR and stop FOMO-only drift.	2-5 days	1-2 weeks	Human email alert renderer + docs alignment.	v0.5.7 PASS
B. Core Dimensions docs/gate	Make full Brevio dimensions unavoidable.	0.5-1 day	2-3 days	Core dimensions doc + gate update.	Docs PR
C. Feedback + Learn/Grow substrate	Natural feedback becomes structured learning.	3-7 days	1-2 weeks	Feedback events, memory contract, tests.	Synthetic feedback smoke
D. Personalized Importance Learning	Reduce noise without going blind.	1-2 weeks	2-4 weeks	Per-user importance memory + evals.	False-positive/negative eval
E. Friend beta 5-15 users	Real private beta.	2-4 weeks	4-8 weeks	More invites, onboarding polish, aftercare.	User feedback + smoke
F. Production FOMO v1	Invite-only product people keep running.	4-7 weeks	8-12 weeks	Stable hosting, monitoring, Google path, support.	Production readiness gate
G. Calendar read/proactivity	Add time/schedule context without writes.	2-4 weeks	4-6 weeks	Calendar read provider + timing signals.	Calendar smoke
H. Personal Context Engine	Broader memory architecture.	3-6 weeks	6-10 weeks	Source-linked user profile.	Memory inspect/delete smoke
I. Drafting assistant	Prepare replies but do not send.	3-6 weeks	6-10 weeks	Draft suggestions + edit/approval loop.	Draft smoke/eval
J. Approved send/write actions	Low/medium-risk actions after exact approval.	4-8 weeks	2-3 months	Email send, calendar writes, confirmations.	Approval-action smoke

K. More providers and MCP tools	Expand beyond Gmail and core tools.	1-3 months	2-4 months	Outlook, tools, adapters.	Provider smoke
L. Browser automation sandbox	Fallback when no API exists.	3-6 months	6-9 months	Sandboxed browser executor.	Site-specific smoke
M. Real-world transactions	Coffee, bookings, travel, purchases with approval.	3-9 months	6-12 months	Prepared actions, spending limits, rollback.	High-risk action smoke
N. Delegated agents + scale	Specialized agents and hundreds/thousands users.	6-12+ months	12-18+ months	Inbox/calendar/memory/travel/safety agents.	SLO/eval gates

7. End-Of-Summer Target

Brevio FOMO private beta with 5-15 controlled users, Personalized Importance Learning started, Human Message Renderer established, friend-safe onboarding, stable SendBlue line decision, Gmail reliability, and basic production hardening.

Stretch target: Brevio FOMO v1 invite-only product that people can actually keep running.

Not realistic by end of summer: full browser automation, payments/bookings/shopping, delegated agents, thousands of users, complete calendar/drafting/sending stack, and full public self-serve compliance at scale.

8. What To Build Next

After the current HMR phase and docs alignment, the next strategic build should be selected by core dimension impact, not by the next FOMO bug.

Likely next serious build candidate: Feedback + Learn/Grow Loop substrate.

Reason: it supports proactivity, Personalized Importance Learning, memory architecture, Human Message Renderer, and future coffee/stocks/travel/calendar behaviors.

9. Claude Operating Rule For All Future Phases

1. Which Brevio dimensions does this phase advance?
2. Which dimensions does it preserve but not advance?
3. Which dimensions are intentionally deferred?
4. Does this phase risk shrinking Brevio into a FOMO/email-only product?
5. Is this only FOMO polish? If yes, why should it not be deferred?
6. What smoke/eval proves the reusable Brevio layer works?

10. Product Philosophy To Preserve

Human Message Renderer

Brevio must turn context into human help. The internal system may reason in metadata, but the user should receive meaning.

Personalized Importance Learning

Brevio must learn each user's personal definition of important. It must become less noisy over time without becoming blind.

Feedback + Learn/Grow Loop

Every interaction should be learnable, but not every interaction should nag the user. Brevio asks more during early training and less as confidence improves.

Safety Gates

The AI may propose; the system gates; the user approves.

11. Final Co-Founder Rule

Do not let Brevio become a polished email alert product. Do not let it become unsafe autonomy either. Build the bridge.

FOMO proves the loop. Brevio expands the loop.

FOMO is the runway. Brevio is the full proactive and autonomous personal agent.

Appendix Update: Composio Strategy For Brevio

This appendix preserves the existing FOMO-to-Brevio map and adds the revised decision: Brevio should use Composio as its integration layer instead of rebuilding Composio, while Brevio keeps control of safety, memory, HMR, approvals, and trust.

12. Composio Adoption Strategy - Integration Layer, Not Brevio Brain

Strategic decision: Brevio should use Composio as the default integration provider instead of rebuilding OAuth plumbing, app connectors, tool schemas, triggers, and provider-specific wrappers. Composio is the integration layer. Brevio remains the product brain, safety system, memory layer, Human Message Renderer, and trust experience.

This updates the tool/workflow orchestration direction: do not build a full internal generic connector platform unless Composio fails a concrete Brevio requirement. Build only the thin Brevio control shell needed to govern Composio safely.

Composio should own	Brevio must own
OAuth/connect links, connected accounts, token handling, provider integrations, tool schemas, tool execution, triggers/webhooks, MCP/direct API compatibility, broad app coverage.	User intent, memory, personalized importance learning, proactivity, permission gates, risk classification, action proposals, audit logs, HMR, user-visible consent and disconnect controls.

Permanent rule: Brevio should not expose raw Composio tools directly to the LLM. Every call must pass through Brevio identity, capability, risk, approval, audit, and HMR layers.

13. Revised Composio Timeline Placement

Stage	Timing	Goal / Deliverable	Hard Gate
0. Adoption decision + fit audit	Now, before any generic internal tool gateway work	Short Composio Fit Report: Google OAuth/testing compatibility, per-user accounts, read/write separation, tool restriction, auditability, revocation, rate limits, failures, first allowed scopes.	No broad code. Decide how to integrate Composio safely, not whether to rebuild Composio.
1. Brevio-Composio foundation	After Google Console F1-F5 and v0.6.0C scope gate, or after active HMR runway if unblocked	composio_user_mapping, connected account status, allowed toolkit/action lists, risk classification, execution wrapper, audit table, feature flag off by default.	No raw tool access to model. No writes. No generic open-ended tool registry.
2. First real connector: Calendar read-only	Next practical integration milestone	Composio-backed Calendar read-only context provider, normalized context, HMR output, one-user smoke, no event create/edit/delete.	Read-only scopes only. Calendar writes remain forbidden until exact approval architecture exists.
3. Missing-access UX	After Feedback + Learn/Grow substrate begins	Brevio can say in plain English when a connected account would help: Connect Calendar, Connect Gmail draft, etc.	User sees capability language, not raw OAuth jargon.
4. Drafting + approved actions	After HMR, feedback learning, and context are stable	Gmail drafts, Gmail send, Calendar writes, Slack posts only through ActionProposal -> risk -> user approval -> Composio execution -> audit -> HMR.	Exact approval required for sends, writes, posts, schedule changes, payments, purchases, bookings.
5. Broader provider expansion	After one or two connectors are proven	Slack, Notion, Drive, GitHub/Linear, CRM as real user needs demand.	Do not add providers because they exist. Add only when they advance a real Brevio dimension or user workflow.

14. What Composio Does Not Change

- FOMO remains the runway. Composio should accelerate reusable Brevio layers, not distract from HMR, feedback, memory, and personalized importance learning.
- The approval law remains unchanged: AI proposes, system gates, user approves.
- High-risk transactions stay deferred: Venmo, Zelle, Stripe transactions, Uber Eats, travel bookings, purchases, and browser automation fallback should not be first-wave Composio integrations.
- Composio success is not measured by number of tools connected. It is measured by whether Brevio can safely understand context, ask for missing access, propose useful help, and execute approved actions with auditability.
- If Composio creates a limitation, Brevio should wrap or constrain it first. Replacing it or building internal equivalents should require proof of a concrete blocker.

Revised one-sentence rule: Adopt Composio early as Brevio's integration provider, but only through a thin Brevio control layer that enforces user identity, allowed tools, read/write separation, risk classification, approval gates, audit logs, and HMR.