

Founder Operating Profile - Brevio

Full proactive/autonomous agent direction update

0. Current Status / How To Use This Profile

This profile is an operating manual for Claude, future agents, engineers, and co-founder-style reviewers. It is not the active phase tracker. Active implementation state lives in FOMO_PLAN.md, apps/fomo/KERNEL.md, current PR reports, and smoke reports.

This updated version re-anchors Brevio around the full agent vision: FOMO is the wedge, but the company goal is a proactive, autonomous, memory-rich, permissioned, human-feeling personal agent.

Founder rule: FOMO should be the runway, not the destination. Brevio is the full agent.

1. Founder Summary

The founder is ambitious, product-obsessed, safety-aware, and impatient with fake progress. He wants a real company-grade AI personal assistant, not a chatbot, not a Gmail script, and not a shallow API wrapper.

He wants Brevio to become an iMessage-native personal AI assistant that feels like a calm helpful person, remembers the user, learns continuously, understands goals, chooses tools, asks for permission, and safely helps across real life.

He wants forward motion, but not reckless motion. Claude should move fast through locked scopes, smoke tests, and small PRs. Claude should challenge weak plans and never reduce Brevio to FOMO-only polish.

2. The Long-Term Product Vision

Text this number. Brevio knows me. It helps me. It asks before doing anything risky.

Brevio should become a proactive personal-agent operating layer. It should learn the user, reason about goals, detect missing access, connect tools with consent, propose help, and eventually complete real-world tasks with approval.

Examples of the long-term product direction:

- Coffee: Brevio may learn the user often orders coffee at a certain time. It may remind, prepare, or ask for approval. It must not silently purchase coffee.
- Stocks: Brevio may learn the user cares about certain stocks, companies, or themes. It may offer watch alerts with permission. It must not silently trade or give reckless financial actions.
- Vacation: Brevio may infer from emails/calendar that the user is planning a trip. It may ask if the user wants help organizing, but it must not silently book or buy anything.

- Calendar: Brevio may learn schedule pressure and suggest timing changes. Writes require permission and exact confirmation.
- Email/drafting: Brevio may draft or suggest, but sending requires explicit user approval.

Brevio can learn patterns. Brevio can propose help. Brevio cannot silently act on high-risk patterns.

3. The Core Brevio Agent Dimensions

Every future phase should state which dimensions it advances, preserves, or defers. This protects the project from becoming a polished FOMO-only product.

Dimension	Founder Meaning
Autonomy	Brevio can reason about what needs to happen, identify missing access, propose plans, and take approved actions.
Proactivity	Brevio notices relevant patterns or events before the user asks, then proposes help safely.
Memory Architecture	Brevio stores source-linked, per-user, correctable, deletable memories with confidence, recency, and audit trails.
Agent Core + Reasoning	Brevio plans, ranks, classifies, decides uncertainty, and chooses the next safe step.
Tool / Workflow Orchestration	Brevio uses tools through governed workflows rather than one-off scripts.
Security / Permission Gates	The AI may propose; the system gates; the user approves.
Multimodal + Perception	Future Brevio can perceive documents, images, screens, files, and context beyond text, only with consent.
Feedback + Learn/Grow Loop	Every interaction should be learnable without nagging the user. Brevio should improve how it serves each person.
Human Message Renderer	Brevio turns structured context into natural, useful, human-feeling messages.
Observability / Evals / Reliability	Smoke tests, audit, evals, recovery, and monitoring make the agent trustworthy.
Production Scale	Hundreds/thousands of users require queues, rate limits, provider contracts, monitoring, support, billing, and compliance.
User Trust / Consent	The user understands what Brevio reads, remembers, proposes, and does; control always remains with the user.

4. FOMO Is The Wedge, Not The Destination

FOMO is valuable because it proves reusable layers: context, ranking, policy gates, Slack approval, SendBlue messaging, reply parsing, memory, audit, and smoke tests. It should not become the company.

FOMO Layer	Long-Term Brevio Layer
Gmail polling	Context provider pattern
Ranker	Reasoning and importance model

Slack approval	Human approval and trust checkpoint
SendBlue text	iMessage-native action channel
STOP/START	Consent, control, and safety semantics
Reply parser	Natural feedback input
Memory signals	Personal memory architecture
Human Message Renderer	Universal product voice layer
Smoke tests	Real-world provider proof
Audit logs	Agent observability

Founder standard: keep FOMO work only when it builds or validates reusable Brevio layers or fixes a real user trust blocker.

5. Permanent Product Layers

5.1 Human Message Renderer

Brevio must turn context into human help. The internal system may reason in metadata, but the user should receive meaning.

Bad: Galiette Mita / Q3 board deck / Time-sensitive sign-off request. Good: Galiette emailed you about the Q3 board deck. It looks time-sensitive - she needs sign-off by tomorrow.

This applies to email, calendar, drafts, tasks, tool results, browser summaries, bookings, payments, memory explanations, and why answers.

5.2 Personalized Importance Learning

Brevio must learn each user's personal definition of important. Urgent wording does not always mean important. Commercial does not always mean unimportant. Learning is per-user, auditable, reversible, and safe.

5.3 Feedback + Learn/Grow Loop

Brevio should learn and grow from user feedback so it can serve the user better over time. Every interaction should be learnable, but not every interaction should nag the user. Brevio should ask more during onboarding and early learning, less as confidence improves.

Brevio's only goal is to serve the user better over time.

6. Architecture And Safety Philosophy

Brevio should become a small operating system for personal AI agents: tools are capabilities, memory/context is the filesystem, permission gates are the kernel, workflows are processes, audit logs are observability, and sandboxed executors isolate dangerous execution.

The eight architecture laws:

1. Tool Lean-in - one clean capability per tool/server/adapter.
2. Context Providers - structured context such as email, calendar, memory, files, contacts, docs.
3. Gateway Connectors - auth, OAuth, consent, routing, rate limits, egress, logging, audit.

4. Stateful Session Managers - tasks, alerts, retries, approvals, workflows, failures.
5. Sandboxed Executors - browser, shell, payment, booking, or computer-use must be isolated.
6. Workflow Packagers - multi-step jobs become named, testable workflows.
7. Delegated Reasoners - future sub-agents have isolated tools/context/permissions/audit.
8. API-first, browser-fallback, approval-required - structured tool first, sandboxed browser fallback only, exact approval before high-risk final actions.

The AI may decide what tools would help. The system decides what tools are allowed. The user approves anything risky.

7. Current Build Strategy

Move fast, but do not let speed turn into FOMO-only churn. After near-term HMR work, the next strategic phases should be selected by core Brevio dimensions, especially Feedback + Learn/Grow, memory architecture, proactivity, and tool orchestration.

Time Horizon	Founder Goal	What To Avoid
Now / next 1-2 weeks	Finish active HMR and core-dimension docs; stop FOMO-only drift.	Do not open random micro-polish phases.
Next 2-4 weeks	Feedback + Learn/Grow substrate; Personalized Importance Learning foundation.	Do not jump to high-risk actions before memory/approval layers.
By end of summer	FOMO private beta with 5-15 controlled users; learning started; stable rails; production hardening.	Do not claim full autonomous Brevio is complete.
After summer	Calendar read, drafting, approved sends, more providers, tool orchestration.	Do not make unsafe silent actions.
6-18 months	Browser fallback, real-world transactions, delegated agents, scale.	Do not skip smoke tests, approval gates, compliance, or memory controls.

8. Claude Operating Rules

- Start every phase with the 6-question gate and a Core Dimension Check.
- State which Brevio dimensions the phase advances, preserves, or defers.
- If the phase only polishes FOMO, explain why it should not be deferred.
- Lock scope before code. No “while I’m here” work.
- Runtime paths must be real or absent. Never half-wired.
- Build/lint/test/PG tests must be reported when relevant.
- Smoke tests prove real integrations before trust.
- No raw private data in logs, fixtures, memory, or repo.
- Do not activate reserved stack items without a real caller.
- Do not reopen OpenAI-vs-Anthropic unless a phase explicitly gates it.

9. Founder Approval Standards

Approve only if the phase moves Brevio forward, preserves safety, and fits the long-term agent OS direction. Pause if Claude overbuilds, underbuilds, repeats work, skips evidence, or treats FOMO as the destination.

The best founder-clone question:

Does this move Brevio closer to a real, trusted, human-feeling personal agent - or does it just make the repo look more advanced?

10. Final Founder Commitments

9. Do not let Claude drift from the full Brevio agent vision.
10. Do not let FOMO become a Gmail script dead end.
11. Do not let infrastructure grow without active callers.
12. Do not skip smoke tests or PASS reports.
13. Do not accept half-wired production paths.
14. Do not weaken privacy or cross-user isolation.
15. Do not let urgent spam destroy trust.
16. Do not let branding, UX, or message rendering be an afterthought.
17. Do not let “move fast” mean “move recklessly.”
18. Do not let Brevio dump metadata as the user experience.
19. Do not let learning be reduced to analytics. Learning is how Brevio gets better at serving the user.

FOMO proves the loop. Brevio expands the loop.

Appendix Update: Founder Operating Rule For Composio

This appendix preserves the existing Founder Operating Profile and adds the revised operating position: use Composio aggressively for integration speed, but never let raw tool access bypass Brevio's control layer.

11. Composio Operating Position

Founder decision: Brevio should use Composio as the actual integration layer. The company should not rebuild a Composio-style connector platform unless a specific, proven requirement makes that unavoidable.

Correct split: Composio is the engine room for integrations. Brevio is the product, brain, memory system, safety layer, user experience, and trust layer.

Use Composio for	Do not delegate to Composio
OAuth/connect flows, connected accounts, token handling, provider APIs, tool schemas, action execution, triggers, MCP/direct API access, broad connector coverage.	What matters to the user, what Brevio remembers, when Brevio should propose help, what is allowed, what is risky, what requires approval, how results are explained, what is audited, what users can correct/delete.

Founder standard: Claude should not build a generic in-house Tool Gateway that duplicates Composio. Claude should build Brevio's thin control shell around Composio: user mapping, capability status, allowlists, read/write separation, risk levels, approvals, audit, and HMR.

12. Composio Implementation Rules For Claude / Agents

- Default posture: use Composio for integrations; do not rebuild provider connectors unless a documented Composio blocker exists.
- Never expose raw Composio tool access directly to the model. All calls go through Brevio's wrapper.
- Every Composio call must have user_id, connected_account_id, tool/action name, risk level, input summary, output summary, status, error category if any, and audit trace.
- Read tools and write tools must be separated. Write tools require ActionProposal and exact user approval unless explicitly classified as safe and pre-approved.
- No Composio provider expansion without a real Brevio use case, smoke test, and Core Dimension Check.
- Feature flags and kill switches are mandatory for Composio execution paths.
- Users must be able to see connected apps, understand what each connection enables, and disconnect/reconnect without engineering intervention.
- Composio does not weaken existing Google OAuth, privacy, cross-user isolation, or approval standards.

13. Revised Build Sequence With Composio

Step	What to build	Why it belongs here
0. Fit audit	Composio Fit Report before internal gateway work.	Locks direction and prevents wasted rebuild work.
1. Foundation	Brevio-Composio user mapping, connected-account status, allowlists, execution wrapper, audit table, feature flag.	Lets Brevio use Composio without making Composio the product.
2. Calendar read-only	First real connector for schedule context only.	Useful for proactivity while avoiding unsafe writes.

3. Missing-access UX	Plain-English connection prompts: connect Calendar, connect Gmail draft, etc.	Makes Brevio easier for nontechnical founders.
4. Approved actions	Draft/send email, create/edit calendar, Slack/Notion actions only through proposal and approval.	Moves toward autonomy without silent risky action.
5. Provider expansion	Add Slack, Notion, Drive, GitHub/Linear, CRM when there is real user value.	Keeps build demand-driven, not connector-list-driven.

Do not start with Venmo, Zelle, Stripe transactions, Uber Eats, purchases, travel bookings, browser automation, or arbitrary any-app action mode. Those require mature approval, audit, spend limits, rollback/failure handling, and user trust controls.

14. Updated Founder Approval Question

When Claude proposes integration work, ask: Is this using Composio to accelerate a real Brevio capability, or is it rebuilding Composio, exposing unsafe tool access, or adding connector theater?

Approve Composio work only when it advances a Brevio dimension, preserves approval gates, and produces real evidence through tests and smoke reports.